

# Resumen Multi-Level Feedback Queue - SistOp

Lautaro Bachmann

## Contents

<b>Scheduling: The Multi-Level Feedback Queue</b>	<b>3</b>
Definitions . . . . .	3
starvation: . . . . .	3
game the scheduler. . . . .	3
voo-doo constants, . . . . .	3
TIP: USE ADVICE WHERE POSSIBLE . . . . .	3
MLFQ: Summary . . . . .	3
why it is called that: . . . . .	3
History is its guide: . . . . .	3
refined set of MLFQ rules, . . . . .	3
MLFQ usefulness . . . . .	4

# Scheduling: The Multi-Level Feedback Queue

## Definitions

### **starvation:**

if there are “too many” interactive jobs in the system, they will combine to consume all CPU time, and thus long-running jobs will never receive any CPU time (they starve).

### **game the scheduler.**

Gaming the scheduler generally refers to the idea of doing something sneaky to trick the scheduler into giving you more than your fair share of the resource.

### **voo-doo constants,**

Constants that seem to require some form of black magic to set them correctly.

## **TIP: USE ADVICE WHERE POSSIBLE**

it is often useful to provide interfaces to allow users or administrators to provide some hints to the OS.

We often call such hints **advice**, as the OS need not necessarily pay attention to it, but rather might take the advice into account in order to make a better decision.

## **MLFQ: Summary**

### **why it is called that:**

it has multiple levels of queues, and uses feedback to determine the priority of a given job.

### **History is its guide:**

Pays attention to how jobs behave over time and treat them accordingly.

### **refined set of MLFQ rules,**

#### **Rule 1:**

If  $\text{Priority}(A) > \text{Priority}(B)$ , A runs (B doesn't).

#### **Rule 2:**

If  $\text{Priority}(A) = \text{Priority}(B)$ , A & B run in round-robin fashion using the time slice (quantum length) of the given queue.

**Rule 3:**

When a job enters the system, it is placed at the highest priority (the topmost queue).

**Rule 4:**

Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue).

**Rule 5:**

After some time period  $S$ , move all the jobs in the system to the topmost queue.

**MLFQ usefulness**

it manages to achieve the best of both worlds: it can deliver excellent overall performance for short-running interactive jobs, and is fair and makes progress for long-running CPU-intensive workloads.